Aiwei Liu

☐ +86 18851830977 | @ liuaw20@mails.tsinghua.edu.cn | ⊕ exlaw.github.io | ⚲ Tsinghua University, Beijing, China

👤 About Me

I am a **final-year Ph.D. candidate** at Tsinghua University, specializing in robust and trustworthy **large language models (LLMs)**. My research has made significant contributions to the field, as evidenced by over 500 citations. Specifically, I focus on three critical areas:

🏷 **LLM Watermarking:** Embedding detectable features in LLM outputs to protect copyright and prevent the misuse of LLMs.

🛡 **LLM Safety Improvement:** Enhancing LLM safety through alignment techniques and red-teaming methods, including adversarial testing, to improve safety against potential vulnerabilities.

🗄 **Robust Natural Language to SQL:** Enhancing the reliability and accuracy of converting natural language queries into SQL, particularly for complex database schemas.

🎓 Education

🏛 **Tsinghua University**   Beijing, China
Ph.D. in Software Engineering; Advisor: Associate Professor Lijie Wen   *Sep 2020 – present*

🏛 **Nanjing University**   Nanjing, China
B.Eng. in Software Engineering; GPA: 4.6/5.0, ranking: 5/220   *Sep 2016 – Jun 2020*

💼 Work Experience

🏛 **University of Illinois Chicago**   Chicago, USA
Visiting Scholar; Advisor: Prof Philip S. Yu (ACM Fellow, IEEE Fellow)   *July 2024 – Dec 2024*
- Project: Privacy of Large Language Models
- Investigated the privacy of watermarked LLMs, specifically their identifiability by users.

🏛 **The Chinese University of Hong Kong**   Hong Kong, China
Visiting Scholar Advisor: Prof Irwin King (ACM Fellow, IEEE Fellow)   *July 2023 – May 2024*
- Project: Watermark for Large Language Models
- 1) Developed an unforgeable publicly verifiable watermark for Large Language Models 2) Write a comprehensive survey about the text watermarking in the era of LLMs.

 Apple AIML Group   Beijing, China
Research Intern: Mentored by Dr. Meng Cao   *Mar 2023 – Sep 2024*
- Project: Prompt Difficulty Evaluation, Safety alignment for LLM
- 1) Developed a LLM-based automatic attributes identification methods for prompt difficulty evaluation. 2) A safety alignment method for LLM that does not require manual annotation of preference data. 3) TIS DPO: Token-level Importance Sampling for Direct Preference Optimization With Estimated Weights

★ First-author Publication Summary

If you are interested in my **first-author** research contributions (maybe important to some institutions in hiring), here is a summary of my **first-author** publications:

📄 **Top-tier Conference Papers:** 4 papers in ICLR (2024, 2025), 1 paper in ACL, 1 paper in KDD, 1 paper in EMNLP, and 1 paper in ACL(Findings) as first author

📓 **Journal Papers:** 1 paper in ACM Computing Surveys (Impact Factor: 23.8) as first author

🏷 SELECTED WORKS IN TEXT WATERMARKING

✎ Can Watermarked LLMs be Identified by Users via Crafted Prompts?
- **Aiwei Liu**, Sheng Guan, Yiming Liu, Leyi Pan, Yifei Zhang, Liancheng Fang, Lijie Wen, Philip S. Yu, Xuming Hu
- **ICLR 2025** (Ranked **3** in all computer science conferences by Google Scholar)

✎ A Semantic invariant Robust Watermark for Large Language Models
- **Aiwei Liu**, Leyi Pan, Xuming Hu, Shiao Meng, Lijie Wen
- **ICLR 2024** (Ranked **3** in all computer science conferences by Google Scholar)

✎ An Unforgeable Publicly Verifiable Watermark for Large Language Models
- **Aiwei Liu**, Leyi Pan, Xuming Hu, Shu'ang Li, Lijie Wen, Irwin King, Philip S. Yu
- **ICLR 2024** (Ranked **3** in all computer science conferences by Google Scholar)

✎ A Survey of Text Watermarking in the Era of Large Language Models
- **Aiwei Liu***, Leyi Pan*, Yijian Lu, Jingjing Li, Xuming Hu, Xi Zhang, Lijie Wen,
  Irwin King, Hui Xiong,Philip S. Yu
- **ACM Computing Surveys** (Impact Factor: **23.8**, ranked **1/143** in Computer Science Theory & Methods)

✎ MarkLLM: An Open-Source Toolkit for LLM Watermarking
- Leyi Pan, **Aiwei Liu**[†] (Project Lead), Zhiwei He, Zitian Gao, Xuandong Zhao, Yijian Lu, Binglin Zhou,
  Shuliang Liu, Xuming Hu, Lijie Wen, Irwin King, Philip S. Yu
- In Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP) (Demo)
- **GitHub repository** has garnered over 300 stars, demonstrating significant community interest and impact.

✎ An Entropy-based Text Watermarking Detection Method
- Yijian Lu, **Aiwei Liu**, Dianzhi Yu, Jingjing Li, Irwin King
- In Proceedings of Association for Computational Linguistics (ACL), 2024

✎ WaterSeeker: Efficient Detection of Watermarked Segments in Large Documents
- Leyi Pan, **Aiwei Liu**, Yijian Lu, Zitian Gao, Yichen Di, Lijie Wen, Irwin King, Philip S. Yu
- In The Findings of North American Chapter of the Association for Computational Linguistics (NAACL), 2025

✎ On the Cross-lingual Consistency of Text Watermark for Large Language Models
- Zhiwei He, Binglin Zhou, Hongkun Hao, **Aiwei Liu**, Xing Wang, Zhaopeng Tu, Zhuosheng Zhang, Rui Wang
- In Proceedings of Association for Computational Linguistics (ACL), 2024 (**Oral**)

🛡 SELECTED WORKS IN LLM SAFETY IMPROVEMENT

✎ TIS-DPO: Token-level Importance Sampling for Direct Preference Optimization With Estimated Weights
- **Aiwei Liu**, Haoping Bai, Zhiyun Lu, Yanchao Sun, Xiang Kong, Xiaoming Wang, Jiulong Shan, Albin Madappally
  Jose, Xiaojiang Liu, Lijie Wen, Philip S. Yu, Meng Cao
- **ICLR 2025** (Ranked **3** in all computer science conferences by Google Scholar)

✎ Direct Large Language Model Alignment Through Self-Rewarding Contrastive Prompt Distillation
- **Aiwei Liu**, Haoping Bai, Zhiyun Lu, Xiang Kong, Simon Wang, Jiulong Shan, Meng Cao, Lijie Wen
- In Proceedings of Association for Computational Linguistics (ACL), 2024

✎ Character-level White-Box Adversarial Attacks against Transformers via Attachable Subwords Substitution
- **Aiwei Liu**, Honghai Yu, Xuming Hu, Shu'ang Li, Li Lin, Fukun Ma, Yawen Yang, Lijie Wen
- In Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2022

🗄 SELECTED WORKS IN ROBUST NATURAL LANGUAGE TO SQL

✎ Semantic Enhanced Text-to-SQL Parsing via Iteratively Learning Schema Linking Graph
- **Aiwei Liu**, Xuming Hu, Li Lin, Lijie Wen
- In Proceedings of the 28th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining
  (SIGKDD), 2022

✎ Exploring the Compositional Generalization in Context Dependent Text-to-SQL Parsing
- **Aiwei Liu**, Wei Liu, Xuming Hu, Shuang Li, Fukun Ma, Yawen Yang, Lijie Wen
- In The Findings of Association for Computational Linguistics (Findings of ACL), 2023

✎ A Comprehensive Evaluation of ChatGPT's Zero-shot Text-to-SQL Capability
- **Aiwei Liu**, Xuming Hu, Lijie Wen, Philip S Yu
- With over **100 citations**, this work has demonstrated significant impact in the field.

🏆 Selected Awards

Tsinghua Excellent Student Award (First Class)	*2024*
Tsinghua Excellent Student Award (SecondClass)	*2022*
Outstanding Graduates Nanjing University	*2020*
China Electronics Technology Group Scholarship	*2019*
National Scholarship	*2018*
Hainan Airlines Scholarship	*2017*

🎤 Professional Services

Program Committee/ Reviewer

- The International Conference on Learning Representations (ICLR)
- The Annual Meeting of the Association for Computational Linguistics (ACL)
- The Annual Conference on Empirical Methods in Natural Language Processing (EMNLP)
- The Annual Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)
- The Annual Conference of the European Chapter of the Association for Computational Linguistics (EACL)
- The ACM WWW International World Wide Web Conference (WWW)
- The ACM International Conference on Multimedia (MM)

Workshop Organization

- Co-organizer of the AAAI 2025 Workshop on Preventing and Detecting LLM Generated Misinformation (PDLM)

📖 Teaching Experience

Preventing and Detecting Misinformation Generated by Large Language Models   July 2024

SIGIR 2024 Tutorial

- **Lead presenter** for a tutorial on techniques for preventing and detecting LLM-generated misinformation at the 47th International ACM SIGIR Conference.
- Tutorial website: **https://sigir24-llm-misinformation.github.io/**

Innovation Talent and University Culture   2021

Teaching Assistant, Tsinghua University

- Assisted in course delivery, facilitated student discussions, and provided support for course-related projects

Operating Systems   2018

Teaching Assistant, Nanjing University

- Conducted lab sessions, graded assignments, and provided one-on-one support to students in understanding complex OS concepts

Other Publications

🔗GDA: Generative Data Augmentation Techniques for Relation Extraction Tasks

- Xuming Hu*, **Aiwei Liu*** (Equal contribution), Zeqi Tan, Xin Zhang, Chenwei Zhang, Irwin King, Philip S. Yu
- In The Findings of Association for Computational Linguistics (Findings of ACL), 2023

🔗RAPL: A Relation-Aware Prototype Learning Approach for Few-Shot Document-Level Relation Extraction

- Shiao Meng, Xuming Hu, **Aiwei Liu**, Shu'ang Li, Fukun Ma, Yawen Yang, Lijie Wen
- In Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP), 2023

🔗AMR-based network for aspect-based sentiment analysisn

- Fukun Ma, Xuming Hu, **Aiwei Liu**, Yawen Yang, Philip S. Yu, Lijie Wen
- In Proceedings of Association for Computational Linguistics (ACL), 2023

🔗CHEF: A Pilot Chinese Dataset for Evidence-Based Fact-Checking

- Xuming Hu, Zhijiang Guo, Guanyu Wu, **Aiwei Liu**, Lijie Wen, Phlip S.Yu

- In The 2022 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL), 2022

⌘A Multi-level Supervised Contrastive Learning Framework for Low-Resource Natural Language Inference
- Shu'ang Li, Xuming Hu, Li Lin, **Aiwei Liu**, Lijie Wen, Philip S. Yu
- IEEE/ACM Transactions on Audio, Speech, and Language Processing

⌘Improving Open Relation Extraction With Search Documents Under Self-Supervisions
- Xuming Hu, Zhaochen Hong, Chenwei Zhang, **Aiwei Liu**, Shiao Meng, Lijie Wen, Irwin King, Philip S. Yu
- IEEE/ACM Transactions on Audio, Speech, and Language Processing

⌘Refiner: Restructure Retrieval Content Efficiently to Advance Question-Answering Capabilities
- Zhonghao Li, Xuming Hu, **Aiwei Liu**, Kening Zheng, Sirui Huang, Hui Xiong
- In The Findings of Empirical Methods in Natural Language Processing (EMNLP), 2024

⌘Entity-to-Text based Data Augmentation for various Named Entity Recognition Tasks
- Xuming Hu, Yong Jiang, **Aiwei Liu**, Zhongqiang Huang, Pengjun Xie, Fei Huang, Lijie Wen, Philip S. Yu
- In The Findings of Association for Computational Linguistics (Findings of ACL), 2023

⌘Enhancing Cross-lingual Natural Language Inference by Soft Prompting with Multilingual Verbalizer
- Shuang Li, Xuming Hu, **Aiwei Liu**, Yawen Yang, Fukun Ma, Philip S. Yu, Lijie Wen
- In The Findings of Association for Computational Linguistics (Findings of ACL), 2023

⌘Gaussian prior reinforcement learning for nested named entity recognition
- Yawen Yang, Xuming Hu, Fukun Ma, Shuang Li, **Aiwei Liu**, Lijie Wen, Philip S. Yu
- IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)

⌘On the Robustness of Document-Level Relation Extraction Models to Entity Name Variations
- Shiao Meng, Xuming Hu, **Aiwei Liu**, Fukun Ma, Yawen Yang, Shuang Li, Lijie Wen
- In The Findings of Association for Computational Linguistics (Findings of ACL), 2024

⌘Entropy-Based Decoding for Retrieval-Augmented Large Language Models
- Zexuan Qiu, Zijing Ou, Bin Wu, Jingjing Li, **Aiwei Liu**, Irwin King
- In The Conference of North American Chapter of the Association for Computational Linguistics (NAACL), 2025

⌘ChatCite: LLM Agent with Human Workflow Guidance for Comparative Literature Summary
- Yutong Li, Lu Chen, **Aiwei Liu**, Kai Yu, Lijie Wen
- Accepted by **Coling 2024**

OTHER PREPRINTS

⌘Interpretable Contrastive Monte Carlo Tree Search Reasoning
- Zitian Gao, Boye Niu, Xuzheng He, Haotian Xu, Hongzhang Liu, **Aiwei Liu**, Xuming Hu, Lijie Wen
- ICLR 2025 Submission

⌘TabGEN-RAG: Iterative Retrieval for Tabular Data Generation with Large Language Models
- Liancheng Fang, **Aiwei Liu**, Hengrui Zhang, Henry Peng Zou, Weizhi Zhang, Philip S. Yu
- Accepted by **TRL@NeurIPS 2024**